

CS 537 - Introduction to Operating Systems - Fall 2017

[Site](#) /

Project3a

Project 3a: Multiprocess Programming with Shared Memory in Linux

You are likely to find that it is not possible to divide this project up by "server" vs "client" (and have one person work on each part) since they need to work together closely; you are likely to have the most success if you work together. This project will not be a lot of code, but it can be tricky to use the suggested library routines correctly. You may find the discussion [here](#) useful.

Objectives

There are three objectives to this assignment:

- To use shared memory across cooperating processes.
- To use mutex for mutual exclusion between processes.
- To catch signals (such as SIGINT) with a signal handler.

Overview

In this assignment, you will implement a client and server that communicate through a shared memory page to display statistics about the client processes. All of the processes are running on the same machine. Each client process will periodically write to the shared memory page with updates about its recent behavior (e.g., how much CPU it has been allocated); the server process collects this information (by reading from the shared memory page) and periodically displays the information for all the client processes. Meanwhile, each client also periodically displays the PIDs of all the processes that are currently running.

As we will soon cover in the lectures, when processes (or threads) cooperate through shared memory, they require synchronization primitives (such as locks or semaphores) to ensure that they do not have race conditions when they are each simultaneously updating the same memory locations. To minimize our need for synchronization in this Project (synchronization is for Project 4!), we will construct the clients and server so that only a single client is (usually) writing to each memory location and the server is only reading (not writing); if the server happens to occasionally

read data that is not up-to-date, that is okay, since the data is just usage statistics (and not your bank account). The only time you will need to worry about mutual exclusion will be when clients are first starting and when they exit.

For this project, you will be implementing two components: a server process that displays client statistics every second, and a client process that sleeps and displays the PIDs of all the processes that are currently running every second.

Server Process

Let us consider what the server process, `shm_server`, must do.

First, it is the responsibility of the server process to create and initialize the shared memory page. A shared memory page can be created with `shm_open()` as follows:

```
int fd = shm_open(SHM_NAME, O_RDWR | O_CREAT, 0660);
```

This routine essentially maps the shared memory page to the `tmpfs` (temporary file system). After calling this routine, you will find this file at `/dev/shm/SHM_NAME`. Therefore, it is very convenient to inspect your shared memory page and debug your code via dumping this file using tools like `xxd`.

Because everyone will create the file at `/dev/shm/`, you have to use a unique file name to avoid conflicts. Therefore, you are required to use the following convention to name this file:

For example, if A works with B for this project, then `SHM_NAME = "aaa_bbb"` where `aaa` is A's CS login and `bbb` is B's CS login. If you are working alone, then simply put your own CS login as the file name.

Second, you should use `ftruncate()` to cause this file to be truncated to a precise size in bytes. For this project, you can only create a single page of shared memory. (Hint: find the page size using some system call)

Third, you should use `mmap()` and the argument flags in this routine should be `MAP_SHARED` because we are sharing this mapping.

If anything goes wrong with this setup (e.g., the shared memory page cannot be exclusively created), then the `shm_server` should exit with return code 1.

After the server process creates this single page of shared memory and initializes it appropriately, it should enter an infinite loop where it continually sleeps for one second, reads the contents of shared memory, and then displays the statistics that have been written to shared memory to STDOUT.

There should be one line of output for each client process; the format of each line of output should be as follows:

```
[Iteration], pid : [pid], birth : [dateOfBirth], elapsed : [sec] s [msec] ms,  
[clientString]
```

The server should have a counter that keeps track of the current iteration and prints it at the start of each line (referring to `Iteration` above).

The `dateOfBirth` should be the string containing the date and time when the client starts in a human-readable format `www Mmm dd hh:mm:ss yyyy`, where `www` is the weekday, `Mmm` the month in letters, `dd` the day of the month, `hh:mm:ss` the time, and `yyyy` the year. You may not need to convert the time to this format yourself. Try to find a C library function to do that for you!

The `elapsed` should be the time that's elapsed since the time of the process started. The `sec` should be an integer and `msec` should be displayed with four digits after the decimal point; you may have either the client or server perform this formatting.

`clientString` is a string that will be initialized for each client when it starts to execute. It will be given by the user as a command line argument to the client.

For example, if there are two client processes, the output may look like the following:

```
4, pid : 14084, birth : Fri Oct 13 22:59:27 2017, elapsed : 30 s 4.8310  
ms, c1  
4, pid : 14085, birth : Fri Oct 13 22:59:27 2017, elapsed : 30 s 4.5600  
ms, c2
```

If at some time, a new client starts writing to the shared memory page, the output could look like this:

```
5, pid : 14084, birth : Fri Oct 13 22:59:27 2017, elapsed : 34 s 5.4670  
ms, c1  
5, pid : 14085, birth : Fri Oct 13 22:59:27 2017, elapsed : 34 s 5.1960  
ms, c2
```

```
5, pid : 14276, birth : Fri Oct 13 23:00:02 2017, elapsed : 0 s 0.0010
ms, c3
```

Finally, if a client (pid : 14085) terminates at some time, the output could look like this:

```
6, pid : 14084, birth : Fri Oct 13 22:59:27 2017, elapsed : 91 s 14.4980
ms, c1
6, pid : 14276, birth : Fri Oct 13 23:00:02 2017, elapsed : 57 s 10.0560
ms, c3
```

You will need to define a structure for each client's statistics. The details of this structure are given below. We suggest treating the shared memory page as an array of these structures. We've provided a definition for the type `stats_t` here.

```
typedef struct {
    int pid;
    char birth[25];
    char clientString[10];
    int elapsed_sec;
    double elapsed_msec;
} stats_t;
```

You must use the fields we have defined without modifying them, but you can add to the structure. However, you cannot add too many fields such that `sizeof(stats_t) > 64`. (See reasonings below.)

To ensure that only one client at a time is searching through and modifying the essential structures of the shared memory page, you should use a mutex to provide mutual exclusion to protect the critical sections.

To use a mutex lock, you should initialize the mutex as:

```
pthread_mutex_init(<mutexVariable>, <mutexAttribute>);
```

Normally, mutexes are process-exclusive, and `<mutexAttribute>` can be set to `NULL` by default. But in this case, a single mutex should be accessed by multiple processes. So, we will have to make some additional changes by initializing `<mutexAttribute>` to be of type `PTHREAD_PROCESS_SHARED`. So, we can do:

```
pthread_mutexattr_init(<mutexAttribute>);  
pthread_mutexattr_setpshared(<mutexAttribute>, PTHREAD_PROCESS_SHARED);
```

And then:

```
pthread_mutex_init(<mutexVariable>, <mutexAttribute>);
```

The clients can then use `pthread_mutex_lock(<mutexVariable>);` to essentially acquire the mutex lock and `pthread_mutex_unlock(<mutexVariable>)` to release the lock.

When your server terminates, you will need to ensure that it correctly removes the shared memory page so that your `/dev/shm/SHM_NAME` would not remain there forever. To do so, you need to call `munmap()` and `shm_unlink()`. Again, read the manuals carefully to learn how to correctly use them!

You need to figure out the maximum number of clients the server should be able to handle (let's call it **maxClients**). If more than **maxClients** clients at a single time try to use this shared memory page, those clients should receive an error.

When a client terminates, it must reset its segment (or some bytes of its segment depending on your implementations) in the shared memory page so that the segment can be used by another client later.

One of your challenges when implementing the server will be to determine the number of valid clients; don't print out any garbage if there are fewer than **maxClients** clients running currently!

Hint: We only created the shared memory for one page. We will break the page into segments, with each segment being of 64 bytes (this is why we limit the size of the structure `stats_t` to be at most 64). The idea is that you will use one segment to store the statistics of one client. In addition, one segment is reserved to store the mutex. You should be able to determine **maxClients** from this information.

Client Process

Your client process, `shm_client`, should behave as follows.

Multiple clients may connect simultaneously to the same server and same shared memory page.

Each client will have 1 command line argument of type string when it starts executing, (which refers to `clientString` above). This string should have length at most 9.

You should first obtain the shared memory page pointer by using `shm_open()` and `mmap()`. Note that there is a subtle difference in the argument of `shm_open()` in the client as opposed to the server because the server has already created the shared memory page and the client should just open this particular preexisting page.

There is no need to initialize the mutex lock in the client (because it was already initialized in the server). All you need to do here is to point the semaphore to the segment that stores the mutex in the shared memory page.

Now your client process will iterate forever. For each iteration, it first updates and writes its statistics to the particular segment of the shared memory page; then it sleeps for one second. Finally, it prints out the PIDs of current active clients. Below is part of the example output from one client:

```
Active clients : 14084 14085 14276
Active clients : 14084 14085 14276
Active clients : 14084 14276
```

We assume that you will kill your client program by sending it a `SIGINT` signal (with Ctrl-C) or a `SIGTERM` signal (by the command line `kill [PID]`). Usually, `SIGINT` and `SIGTERM` interrupts your program and kills it. However, you can change this default behavior by specifying a signal handler that should be run when that particular signal is delivered. To do this, use `sigaction()` to specify the routine that you want to run. This new routine should make sure that the segment is reset (or some bytes of its segment depending on your implementations) such that the associated segment can be used by another client later, and then call `exit`. You may find online examples like [this](#) very useful.

Similarly, the server will be killed by receiving the signal. So clean the shared memory page and exit elegantly when the server terminates.

To help you start , we have provided the [code frame of shm_server.c and shm_client.c](#).

Compiling, Makefile, and Testing

You should compile your code like

```
gcc shm_server.c -o shm_server -Wall -Werror -lrt -lpthread
```

To ensure that we compile your server and your client process correctly, you will need to create a simple Makefile. The makefile must make both of the targets. If you don't know how to write a

makefile, you might want to look at the man pages for make.

To run the test scripts, go to the directory where your source codes reside and type

```
~cs537-1/ta/tests/3a/runtests -c
```

Page last modified on October 18, 2017, at 11:15 AM

